

AI-risk coverage and AI equities: an exploratory event study

Observed window: September 8-30, 2026
Prepared: September 30, 2026

Executive summary

On September 8, Anthropic employee Evan Hubinger replied on X that AI could kill all humans and put his personal probability above 10% within a decade, following Jacob Coxon’s resignation warning. [Axios](#), [BBC](#), and [Ars Technica](#) documented the exchange.

The reproducible news sample is sharply risk-framed: mean VADER is **−0.400** versus **+0.014** for broad AI, and risk intensity is **0.158** versus **0.010**. Across 16 trading observations, the equal-weight AI basket returned **+2.88%** versus **−0.19%** for SPY; news-tone Granger tests do not reject no predictive content ($p = .685$ and $.961$).

The relevance-ranked X-post sample is instead slightly more positive and less risk-intense than broad-AI posts. Neither source has a significant one-lag Granger result; these 16 observations do not support a causal or return-prediction claim.

Event, scope, and data

The event date is September 8, 2026, based on the supplied X post. The empirical window ends September 30, 2026; all tables, figures, and diagnostics were rerun together after that close.

Social and news evidence

The supplied X posts anchor the event, not a representative X sample. The news notebook caches two dated Google News RSS searches for September 8–30.

- 1. **Event/risk sample:** Hubinger, Coxon, or Anthropic combined with AI-risk/existential/humanity terms. It yielded 56 headlines on eight calendar dates (September 9, 10, 11, 13, 14, 16, 23, and 29).
- 2. **Broad-AI comparison sample:** “artificial intelligence” or AI. It yielded 99 headlines in the same window.

Headlines receive VADER compound sentiment and risk intensity, **(negative-risk terms – positive/safety terms) / words**. This adapts the COVID-study comparison logic to a small, provider-ranked RSS sample; it cannot identify general-media bias.

The original notebook displays the following first ten dated event-query headlines; this is an exact notebook output, not a separate dataset.

Date	Headline as displayed in notebook	Source
Sep. 9	Ex-Anthropic Researcher Warns AI Threat to Hum...	조선일보
Sep. 9	A researcher warned AI could end humanity. Con...	USA Today

Date	Headline as displayed in notebook	Source
Sep. 9	Anthropic Researchers Raise Alarm Over A.I. Ac...	The New York Times
Sep. 9	Anthropic insiders warn AI could kill all huma...	Axios
Sep. 9	Experts weigh in as researcher says AI has mor...	CNBC
Sep. 9	Anthropic Worker Quits Over AI Firms 'Gambling...	bloomberg.com
Sep. 9	AI could kill humanity, warns Anthropic scient...	Firstpost
Sep. 9	Anthropic researcher believes more than 10% ch...	BBC
Sep. 9	AI researchers 'earnestly believe' it could ki...	CBC
Sep. 9	Could AI wipe out humanity in a decade? An Ant...	The Independent

The attached Twitter and finance studies motivate daily aggregation, log returns, and transparent sentiment proxies; short-text classification remains a limitation.

X-post archive sample

[TwitterAPIio_Archive_Sentiment_Analysis.ipynb](#) caches relevance-ranked Advanced Search results, applies the same measures, and filters New York timestamps to September 8–30. The samples span 22 calendar days but are query-defined rather than an X census.


X-post comparison	Event/risk query	Broad-AI query	Difference
Posts after date filter	1,086	1,722	−636
Mean VADER compound	+0.1640	+0.1435	+0.0205
Mean risk intensity	−0.0074	+0.0007	−0.0080

The targeted X sample is modestly **more positive** and has **lower** lexicon risk intensity than the broad-AI sample. It therefore supplies no analogous evidence of negative AI-risk coverage bias; the contrast with the news result illustrates source-specific framing, not platform-wide sentiment.

One-lag X-post tests have p-values **.2010**, **.9117**, and **.8876**; none rejects at 5%.

Lag-1 X-post driver	p-value	Usable observations	Result at 5%
Mean compound sentiment	0.2010	16	Do not reject
Risk intensity	0.9117	16	Do not reject
Post count	0.8876	16	Do not reject

X-post volume peaks at 104 on September 9 and 235 on September 30; tone is otherwise mixed. These are ranking-dependent sample patterns, and mixed excess returns plus the non-significant tests provide no reliable leading-return relationship.



The full executed X-post panel is below. As in the notebook, returns are decimal log returns and 0.0000 on September 23 represents no returned targeted post, not neutral social sentiment.

Date	Basket return	Basket excess return	Posts	Mean compound	Risk intensity	Engagement
Sep. 9	-0.0016	0.0031	104	-0.0834	-0.0152	138056
Sep. 10	-0.0084	-0.0024	33	-0.2949	-0.0020	41770
Sep. 11	0.0091	0.0006	17	0.0494	0.0003	1614
Sep. 14	-0.0077	-0.0032	19	0.0732	-0.0076	26260
Sep. 15	-0.0090	-0.0044	24	0.2585	-0.0095	56740
Sep. 16	-0.0006	0.0038	3	0.1505	-0.0097	67
Sep. 17	0.0200	0.0087	15	-0.0752	-0.0034	24735
Sep. 18	0.0052	0.0039	6	0.3364	-0.0114	8490
Sep. 21	0.0315	0.0161	4	0.3087	-0.0051	10584
Sep. 22	-0.0015	-0.0014	4	-0.2570	-0.0106	1101
Sep. 23	-0.0142	-0.0070	0	0.0000	0.0000	0

Date	Basket return	Basket excess return	Posts	Mean compound	Risk intensity	Engagement
Sep. 24	0.0065	0.0073	23	0.4566	−0.0187	11861
Sep. 25	0.0022	−0.0032	80	0.2212	−0.0089	2694
Sep. 28	−0.0097	−0.0022	109	0.3743	−0.0122	6134
Sep. 29	0.0065	0.0083	181	0.0679	0.0078	7179
Sep. 30	0.0001	0.0022	235	0.2802	−0.0042	6088

AI basket and returns

The basket contains NVIDIA (NVDA), Meta (META), Amazon (AMZN), Alphabet (GOOGL), Microsoft (MSFT), Broadcom (AVGO), and Taiwan Semiconductor (TSM). These firms span AI accelerators/foundry capacity, hyperscale cloud/model deployment, software platforms, and networking/custom silicon. It is an **equal-weight** basket: this is a transparent exposure screen, not a recommendation or a market-cap index. The benchmark is SPY. Yahoo Finance adjusted closes are downloaded by [yfinance](#); the notebook calculates daily log return, the equal-weight basket return, and basket excess return ([basket](#) − [SPY](#)).


Results

Trend in coverage tone

The event/risk sample is substantially more negative than the broad-AI sample.

Measure	Event/risk query	Broad-AI query	Difference
Headlines	56.000000	99.000000	—
Mean VADER compound	−0.399796	+0.014068	−0.413864
Mean risk intensity	0.157827	0.010165	+0.147662

Coverage is episodic: 10 headlines on September 9 (mean −0.516), negative covered days through September 23, and 31 headlines on September 29. The September 29 IPO-risk story is a later shock, not mechanically the September 8 event. Zeros mean no sampled headline, not neutral sentiment.



The chart and full aligned daily panel below reproduce the original notebook's executed outputs. Returns are decimal log returns; `0.0000` for the news variables denotes no sampled event-query headline, not neutral public sentiment.

Date	Headlines	Mean compound	Risk intensity	Basket return	Basket excess return
Sep. 9	10	-0.5160	0.1096	-0.0016	0.0031
Sep. 10	2	-0.7211	0.1083	-0.0084	-0.0024
Sep. 11	3	-0.5356	0.0778	0.0091	0.0006
Sep. 14	4	-0.4070	0.1234	-0.0077	-0.0032
Sep. 15	0	0.0000	0.0000	-0.0090	-0.0044
Sep. 16	3	-0.4607	0.1590	-0.0006	0.0038
Sep. 17	0	0.0000	0.0000	0.0200	0.0087
Sep. 18	0	0.0000	0.0000	0.0052	0.0039
Sep. 21	0	0.0000	0.0000	0.0315	0.0161
Sep. 22	0	0.0000	0.0000	-0.0015	-0.0014
Sep. 23	2	-0.4450	0.0911	-0.0142	-0.0070
Sep. 24	0	0.0000	0.0000	0.0065	0.0073
Sep. 25	0	0.0000	0.0000	0.0022	-0.0032
Sep. 28	0	0.0000	0.0000	-0.0097	-0.0022
Sep. 29	31	-0.3316	0.1981	0.0065	0.0083
Sep. 30	0	0.0000	0.0000	0.0001	0.0022

Within this event-focused RSS sample, risk framing is much more negative than broad AI coverage. Search selection and RSS ranking prevent a general-media-bias inference.

Equity performance

Across September 9–30, the equal-weight basket compounded **+2.88%** versus **−0.19%** for SPY. Mixed daily excess returns do not show a stable relation to negative coverage. This is a descriptive comparison, not abnormal-return alpha.

Granger diagnostic

For each driver, the notebook runs a one-lag bivariate Granger test with basket excess log return as the dependent series. The null is that the lagged driver does not improve the return forecast.

Lag-1 driver	p-value	Result at 5%
Mean headline compound	0.6851	Do not reject
Risk intensity	0.9607	Do not reject
Headline count	0.8533	Do not reject

With only 16 observations, these bivariate tests are low power and non-causal; they neither establish nor rule out structural effects. A stronger design needs a longer window, richer text, independent social sampling, and market-model controls.

Conclusion

The September 8 warning produced conspicuously negative sampled news coverage, but not negative relative tone in the relevance-ranked X-post sample. The news difference is a measurable **risk-framing differential**; the cross-source contrast shows why it is not evidence of general-media or platform-wide bias.

The basket outperformed SPY, but neither source supplies a significant next-day predictive relationship. The short sample supports no causal claim.

Sources and methodological references

- Hubinger/Coxon event reporting: [Axios](#); [BBC](#); [Ars Technica](#).
- X-post collection: [TwitterAPI.io Advanced Search](#).
- Price source: [Yahoo Finance historical data](#).
- News discovery source: [Google News RSS](#); the notebook records the exact dated query and caches the response for audit.
- Carvalho, J. & Plastino, A. (2020). *On the evaluation and combination of state-of-the-art features in Twitter sentiment analysis*. Artificial Intelligence Review.
- Hogenboom, A. et al. (2014). *Twitter sentiment analysis applied to finance: A case study in the retail industry*. KDIR.
- Sacerdote, B., Sehgal, R. & Cook, M. (2020). *Why Is All COVID News Bad News?* NBER Working Paper 28110.